

# Bits, Bytes, and Integers

15-213: Introduction to Computer Systems  
2<sup>nd</sup> and 3<sup>rd</sup> Lectures, Sep. 3 and Sep. 8, 2015

**Instructors:**


Randal E. Bryant and David R. O'Hallaron

# Today: Bits, Bytes, and Integers

- **Representing information as bits**
- **Bit-level manipulations**
- **Integers**
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting
  - Summary
- **Representations in memory, pointers, strings**

# Everything is bits

- Each bit is 0 or 1
- By encoding/interpreting sets of bits in various ways
  - Computers determine what to do (instructions)
  - ... and represent and manipulate numbers, sets, strings, etc...
- Why bits? Electronic Implementation
  - Easy to store with bistable elements
  - Reliably transmitted on noisy and inaccurate wires



# For example, can count in binary

## ■ Base 2 Number Representation

- Represent  $15213_{10}$  as  $11101101101101_2$
- Represent  $1.20_{10}$  as  $1.0011001100110011[0011]..._2$
- Represent  $1.5213 \times 10^4$  as  $1.1101101101101_2 \times 2^{13}$

# Encoding Byte Values

## ■ Byte = 8 bits

- Binary  $00000000_2$  to  $11111111_2$
- Decimal:  $0_{10}$  to  $255_{10}$
- Hexadecimal  $00_{16}$  to  $FF_{16}$ 
  - Base 16 number representation
  - Use characters '0' to '9' and 'A' to 'F'
  - Write  $FA1D37B_{16}$  in C as
    - `0xFA1D37B`
    - `0xfa1d37b`

| Hex | Decimal | Binary |
|-----|---------|--------|
| 0   | 0       | 0000   |
| 1   | 1       | 0001   |
| 2   | 2       | 0010   |
| 3   | 3       | 0011   |
| 4   | 4       | 0100   |
| 5   | 5       | 0101   |
| 6   | 6       | 0110   |
| 7   | 7       | 0111   |
| 8   | 8       | 1000   |
| 9   | 9       | 1001   |
| A   | 10      | 1010   |
| B   | 11      | 1011   |
| C   | 12      | 1100   |
| D   | 13      | 1101   |
| E   | 14      | 1110   |
| F   | 15      | 1111   |

# Example Data Representations

| C Data Type              | Typical 32-bit | Typical 64-bit | x86-64 |
|--------------------------|----------------|----------------|--------|
| <code>char</code>        | 1              | 1              | 1      |
| <code>short</code>       | 2              | 2              | 2      |
| <code>int</code>         | 4              | 4              | 4      |
| <code>long</code>        | 4              | 8              | 8      |
| <code>float</code>       | 4              | 4              | 4      |
| <code>double</code>      | 8              | 8              | 8      |
| <code>long double</code> | –              | –              | 10/16  |
| <code>pointer</code>     | 4              | 8              | 8      |

# Today: Bits, Bytes, and Integers

- Representing information as bits
- **Bit-level manipulations**
- **Integers**
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting
  - Summary
- Representations in memory, pointers, strings

# Boolean Algebra

## ■ Developed by George Boole in 19th Century

- Algebraic representation of logic
  - Encode “True” as 1 and “False” as 0

### And

- $A \& B = 1$  when both  $A=1$  and  $B=1$

| $\&$ | 0 | 1 |
|------|---|---|
| 0    | 0 | 0 |
| 1    | 0 | 1 |

### Or

- $A | B = 1$  when either  $A=1$  or  $B=1$

| $ $ | 0 | 1 |
|-----|---|---|
| 0   | 0 | 1 |
| 1   | 1 | 1 |

### Not

- $\sim A = 1$  when  $A=0$

| $\sim$ |   |
|--------|---|
| 0      | 1 |
| 1      | 0 |

### Exclusive-Or (Xor)

- $A \wedge B = 1$  when either  $A=1$  or  $B=1$ , but not both

| $\wedge$ | 0 | 1 |
|----------|---|---|
| 0        | 0 | 1 |
| 1        | 1 | 0 |



# General Boolean Algebras

## ■ Operate on Bit Vectors

- Operations applied bitwise

|                       |                   |                   |                   |
|-----------------------|-------------------|-------------------|-------------------|
| 01101001              | 01101001          | 01101001          | 01101001          |
| <u>&amp; 01010101</u> | <u>  01010101</u> | <u>^ 01010101</u> | <u>~ 01010101</u> |
| 01000001              | 01111101          | 00111100          | 10101010          |

## ■ All of the Properties of Boolean Algebra Apply

# Example: Representing & Manipulating Sets

## ■ Representation

- Width  $w$  bit vector represents subsets of  $\{0, \dots, w-1\}$

- $a_j = 1$  if  $j \in A$

- 01101001       $\{0, 3, 5, 6\}$

- 76543210

- 01010101       $\{0, 2, 4, 6\}$

- 76543210

## ■ Operations

- & Intersection      01000001       $\{0, 6\}$
- | Union      01111101       $\{0, 2, 3, 4, 5, 6\}$
- ^ Symmetric difference      00111100       $\{2, 3, 4, 5\}$
- ~ Complement      10101010       $\{1, 3, 5, 7\}$

# Bit-Level Operations in C

## ■ Operations &, |, ~, ^ Available in C

- Apply to any “integral” data type
  - long, int, short, char, unsigned
- View arguments as bit vectors
- Arguments applied bit-wise

## ■ Examples (Char data type)

- $\sim 0x41 \rightsquigarrow 0xBE$ 
  - $\sim 01000001_2 \rightsquigarrow 10111110_2$
- $\sim 0x00 \rightsquigarrow 0xFF$ 
  - $\sim 00000000_2 \rightsquigarrow 11111111_2$
- $0x69 \& 0x55 \rightsquigarrow 0x41$ 
  - $01101001_2 \& 01010101_2 \rightsquigarrow 01000001_2$
- $0x69 | 0x55 \rightsquigarrow 0x7D$ 
  - $01101001_2 | 01010101_2 \rightsquigarrow 01111101_2$

# Contrast: Logic Operations in C

## ■ Contrast to Logical Operators

- `&&`, `||`, `!`
  - View 0 as “False”
  - Anything nonzero as “True”
  - Always return 0 or 1
  - Early termination

## ■ Examples (char data type)

- `!0x41`  $\leadsto$  `0x00`
- `!0x00`  $\leadsto$  `0x01`
- `!!0x41`  $\leadsto$  `0x01`
- `0x69 && 0x55`  $\leadsto$  `0x01`
- `0x69 || 0x55`  $\leadsto$  `0x01`
- `p && *p` (avoids null pointer access)

# Contrast: Logic Operations in C

## ■ Contrast to Logical Operators

- `&&`, `||`, `!`
  - View 0 as “False”
  - Anything non-zero is “True”
  - Always evaluates both sides
  - Early exit

## ■ Example

- `!0x41`  $\leadsto$  0
- `!0x00`  $\leadsto$  1
- `!!0x41`  $\leadsto$  1
- `0x69 && 0x55`  $\leadsto$  0x01
- `0x69 || 0x55`  $\leadsto$  0x01
- `p && *p` (avoids null pointer access)

**Watch out for `&&` vs. `&` (and `||` vs. `|`)...  
one of the more common oopsies in  
C programming**

# Shift Operations

- **Left Shift:  $x \ll y$** 
  - Shift bit-vector  $x$  left  $y$  positions
    - Throw away extra bits on left
    - Fill with 0's on right
- **Right Shift:  $x \gg y$** 
  - Shift bit-vector  $x$  right  $y$  positions
    - Throw away extra bits on right
  - Logical shift
    - Fill with 0's on left
  - Arithmetic shift
    - Replicate most significant bit on left
- **Undefined Behavior**
  - Shift amount  $< 0$  or  $\geq$  word size

|                |          |
|----------------|----------|
| Argument $x$   | 01100010 |
| $\ll 3$        | 00010000 |
| Log. $\gg 2$   | 00011000 |
| Arith. $\gg 2$ | 00011000 |

|                |          |
|----------------|----------|
| Argument $x$   | 10100010 |
| $\ll 3$        | 00010000 |
| Log. $\gg 2$   | 00101000 |
| Arith. $\gg 2$ | 11101000 |

# Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
  - **Representation: unsigned and signed**
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting
  - Summary
- Representations in memory, pointers, strings
- Summary

# Encoding Integers

## Unsigned

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

## Two's Complement

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

```
short int x = 15213;
short int y = -15213;
```

Sign  
Bit



### ■ C short 2 bytes long

|          | Decimal | Hex   | Binary            |
|----------|---------|-------|-------------------|
| <b>x</b> | 15213   | 3B 6D | 00111011 01101101 |
| <b>y</b> | -15213  | C4 93 | 11000100 10010011 |

### ■ Sign Bit

- For 2's complement, most significant bit indicates sign
  - 0 for nonnegative
  - 1 for negative



## Two-complement Encoding Example (Cont.)

**x =**        15213: 00111011 01101101  
**y =**        -15213: 11000100 10010011

| Weight     | 15213        |      | -15213        |        |
|------------|--------------|------|---------------|--------|
| 1          | 1            | 1    | 1             | 1      |
| 2          | 0            | 0    | 1             | 2      |
| 4          | 1            | 4    | 0             | 0      |
| 8          | 1            | 8    | 0             | 0      |
| 16         | 0            | 0    | 1             | 16     |
| 32         | 1            | 32   | 0             | 0      |
| 64         | 1            | 64   | 0             | 0      |
| 128        | 0            | 0    | 1             | 128    |
| 256        | 1            | 256  | 0             | 0      |
| 512        | 1            | 512  | 0             | 0      |
| 1024       | 0            | 0    | 1             | 1024   |
| 2048       | 1            | 2048 | 0             | 0      |
| 4096       | 1            | 4096 | 0             | 0      |
| 8192       | 1            | 8192 | 0             | 0      |
| 16384      | 0            | 0    | 1             | 16384  |
| -32768     | 0            | 0    | 1             | -32768 |
| <b>Sum</b> | <b>15213</b> |      | <b>-15213</b> |        |

# Numeric Ranges

## ■ Unsigned Values

- $UMin = 0$   
000...0
- $UMax = 2^w - 1$   
111...1

## ■ Two's Complement Values

- $TMin = -2^{w-1}$   
100...0
- $TMax = 2^{w-1} - 1$   
011...1

## ■ Other Values

- Minus 1  
111...1

### Values for $W = 16$

|             | Decimal       | Hex          | Binary            |
|-------------|---------------|--------------|-------------------|
| <b>UMax</b> | <b>65535</b>  | <b>FF FF</b> | 11111111 11111111 |
| <b>TMax</b> | <b>32767</b>  | <b>7F FF</b> | 01111111 11111111 |
| <b>TMin</b> | <b>-32768</b> | <b>80 00</b> | 10000000 00000000 |
| <b>-1</b>   | <b>-1</b>     | <b>FF FF</b> | 11111111 11111111 |
| <b>0</b>    | <b>0</b>      | <b>00 00</b> | 00000000 00000000 |

# Values for Different Word Sizes

|      | W    |         |                |                            |
|------|------|---------|----------------|----------------------------|
|      | 8    | 16      | 32             | 64                         |
| UMax | 255  | 65,535  | 4,294,967,295  | 18,446,744,073,709,551,615 |
| TMax | 127  | 32,767  | 2,147,483,647  | 9,223,372,036,854,775,807  |
| TMin | -128 | -32,768 | -2,147,483,648 | -9,223,372,036,854,775,808 |

## ■ Observations

- $|TMin| = TMax + 1$ 
  - Asymmetric range
- $UMax = 2 * TMax + 1$

## ■ C Programming

- `#include <limits.h>`
- Declares constants, e.g.,
  - `ULONG_MAX`
  - `LONG_MAX`
  - `LONG_MIN`
- Values platform specific

# Unsigned & Signed Numeric Values

| X    | B2U(X) | B2T(X) |
|------|--------|--------|
| 0000 | 0      | 0      |
| 0001 | 1      | 1      |
| 0010 | 2      | 2      |
| 0011 | 3      | 3      |
| 0100 | 4      | 4      |
| 0101 | 5      | 5      |
| 0110 | 6      | 6      |
| 0111 | 7      | 7      |
| 1000 | 8      | -8     |
| 1001 | 9      | -7     |
| 1010 | 10     | -6     |
| 1011 | 11     | -5     |
| 1100 | 12     | -4     |
| 1101 | 13     | -3     |
| 1110 | 14     | -2     |
| 1111 | 15     | -1     |

## ■ Equivalence

- Same encodings for nonnegative values

## ■ Uniqueness

- Every bit pattern represents unique integer value
- Each representable integer has unique bit encoding

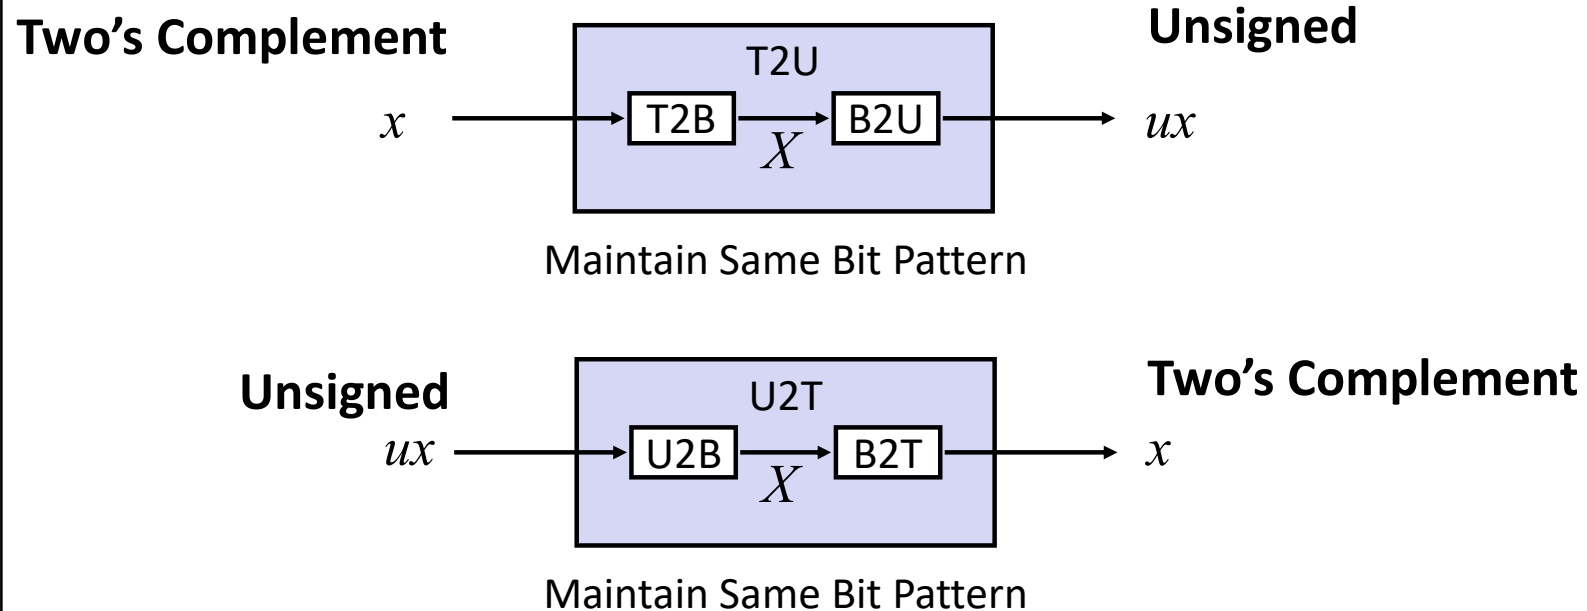
## ■ $\Rightarrow$ Can Invert Mappings

- $U2B(x) = B2U^{-1}(x)$ 
  - Bit pattern for unsigned integer
- $T2B(x) = B2T^{-1}(x)$ 
  - Bit pattern for two's comp integer

# Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
  - Representation: unsigned and signed
  - **Conversion, casting**
  - Expanding, truncating
  - Addition, negation, multiplication, shifting
  - Summary
- Representations in memory, pointers, strings

# Mapping Between Signed & Unsigned



- Mappings between unsigned and two's complement numbers:  
**Keep bit representations and reinterpret**

# Mapping Signed $\leftrightarrow$ Unsigned

| Bits | Signed |         | Unsigned |
|------|--------|---------|----------|
| 0000 | 0      |         | 0        |
| 0001 | 1      |         | 1        |
| 0010 | 2      |         | 2        |
| 0011 | 3      |         | 3        |
| 0100 | 4      |         | 4        |
| 0101 | 5      |         | 5        |
| 0110 | 6      |         | 6        |
| 0111 | 7      |         | 7        |
| 1000 | -8     | → T2U → | 8        |
| 1001 | -7     | ← U2T ← | 9        |
| 1010 | -6     |         | 10       |
| 1011 | -5     |         | 11       |
| 1100 | -4     |         | 12       |
| 1101 | -3     |         | 13       |
| 1110 | -2     |         | 14       |
| 1111 | -1     |         | 15       |



# Mapping Signed $\leftrightarrow$ Unsigned

| Bits | Signed |                                 | Unsigned |
|------|--------|---------------------------------|----------|
| 0000 | 0      | $\longleftrightarrow$<br>=      | 0        |
| 0001 | 1      |                                 | 1        |
| 0010 | 2      |                                 | 2        |
| 0011 | 3      |                                 | 3        |
| 0100 | 4      |                                 | 4        |
| 0101 | 5      |                                 | 5        |
| 0110 | 6      |                                 | 6        |
| 0111 | 7      |                                 | 7        |
| 1000 | -8     | $\longleftrightarrow$<br>+/- 16 | 8        |
| 1001 | -7     |                                 | 9        |
| 1010 | -6     |                                 | 10       |
| 1011 | -5     |                                 | 11       |
| 1100 | -4     |                                 | 12       |
| 1101 | -3     |                                 | 13       |
| 1110 | -2     |                                 | 14       |
| 1111 | -1     |                                 | 15       |



# Relation between Signed & Unsigned

Two's Complement




**Large negative weight**  
*becomes*  
**Large positive weight**

# Conversion Visualized

## ■ 2's Comp. → Unsigned

- Ordering Inversion
- Negative → Big Positive



# Signed vs. Unsigned in C

## ■ Constants

- By default are considered to be signed integers
- Unsigned if have “U” as suffix

0U, 4294967259U

## ■ Casting

- Explicit casting between signed & unsigned same as U2T and T2U

```
int tx, ty;  
unsigned ux, uy;  
tx = (int) ux;  
uy = (unsigned) ty;
```

- Implicit casting also occurs via assignments and procedure calls

```
tx = ux;  
uy = ty;
```

# Casting Surprises

## ■ Expression Evaluation

- If there is a mix of unsigned and signed in single expression,  
*signed values implicitly cast to unsigned*
- Including comparison operations  $<$ ,  $>$ ,  $==$ ,  $<=$ ,  $>=$
- Examples for  $W = 32$ : **TMIN = -2,147,483,648**, **TMAX = 2,147,483,647**

| ■ Constant <sub>1</sub> | Constant <sub>2</sub> | Relation | Evaluation |
|-------------------------|-----------------------|----------|------------|
| 0                       | 0U                    | ==       | unsigned   |
| -1                      | 0                     | <        | signed     |
| -1                      | 0U                    | >        | unsigned   |
| 2147483647              | -2147483647-1         | >        | signed     |
| 2147483647U             | -2147483647-1         | <        | unsigned   |
| -1                      | -2                    | >        | signed     |
| (unsigned)-1            | -2                    | >        | unsigned   |
| 2147483647              | 2147483648U           | <        | unsigned   |
| 2147483647              | (int) 2147483648U     | >        | signed     |

## Summary

### Casting Signed $\leftrightarrow$ Unsigned: Basic Rules

- Bit pattern is maintained
- But reinterpreted
- Can have unexpected effects: adding or subtracting  $2^w$
- Expression containing signed and unsigned int
  - `int` is cast to `unsigned`!!

# Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
  - Representation: unsigned and signed
  - Conversion, casting
  - **Expanding, truncating**
  - Addition, negation, multiplication, shifting
  - Summary
- Representations in memory, pointers, strings


# Sign Extension

## ■ Task:

- Given  $w$ -bit signed integer  $x$
- Convert it to  $w+k$ -bit integer with same value

## ■ Rule:

- Make  $k$  copies of sign bit:
- $X' = \underbrace{x_{w-1}, \dots, x_{w-1}}_{k \text{ copies of MSB}}, x_{w-1}, x_{w-2}, \dots, x_0$



# Sign Extension Example

```
short int x = 15213;
int      ix = (int) x;
short int y = -15213;
int      iy = (int) y;
```

|           | Decimal | Hex         | Binary                              |
|-----------|---------|-------------|-------------------------------------|
| <b>x</b>  | 15213   | 3B 6D       | 00111011 01101101                   |
| <b>ix</b> | 15213   | 00 00 3B 6D | 00000000 00000000 00111011 01101101 |
| <b>y</b>  | -15213  | C4 93       | 11000100 10010011                   |
| <b>iy</b> | -15213  | FF FF C4 93 | 11111111 11111111 11000100 10010011 |

- Converting from smaller to larger integer data type
- C automatically performs sign extension



# Summary:

## Expanding, Truncating: Basic Rules

- **Expanding (e.g., short int to int)**
  - Unsigned: zeros added
  - Signed: sign extension
  - Both yield expected result
  
- **Truncating (e.g., unsigned to unsigned short)**
  - Unsigned/signed: bits are truncated
  - Result reinterpreted
  - Unsigned: mod operation
  - Signed: similar to mod
  - For small numbers yields expected behavior

# Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - **Addition, negation, multiplication, shifting**
- Representations in memory, pointers, strings
- Summary

# Unsigned Addition

Operands:  $w$  bits


$u$  

$+ v$  

True Sum:  $w+1$  bits

$u + v$  

Discard Carry:  $w$  bits

$\text{UAdd}_w(u, v)$  

## ■ Standard Addition Function

- Ignores carry output


## ■ Implements Modular Arithmetic

$$s = \text{UAdd}_w(u, v) = u + v \bmod 2^w$$

# Visualizing (Mathematical) Integer Addition

## ■ Integer Addition

- 4-bit integers  $u, v$
- Compute true sum  $\text{Add}_4(u, v)$
- Values increase linearly with  $u$  and  $v$
- Forms planar surface




# Visualizing Unsigned Addition


## ■ Wraps Around

- If true sum  $\geq 2^w$
- At most once

True Sum



Modular Sum



# Two's Complement Addition

Operands:  $w$  bits


$u$  

$+$   $v$  

True Sum:  $w+1$  bits

$u + v$  

Discard Carry:  $w$  bits

$\text{TAdd}_w(u, v)$  

## ■ TAdd and UAdd have Identical Bit-Level Behavior

- Signed vs. unsigned addition in C:

```
int s, t, u, v;
```

```
s = (int) ((unsigned) u + (unsigned) v);
```


```
t = u + v
```

- Will give `s == t`

# TAdd Overflow

## ■ Functionality

- True sum requires  $w+1$  bits
- Drop off MSB
- Treat remaining bits as 2's comp. integer



# Visualizing 2's Complement Addition


## ■ Values

- 4-bit two's comp.
- Range from -8 to +7

## ■ Wraps Around

- If  $\text{sum} \geq 2^{w-1}$ 
  - Becomes negative
  - At most once
- If  $\text{sum} < -2^{w-1}$ 
  - Becomes positive
  - At most once

NegOver






# Multiplication

- **Goal: Computing Product of  $w$ -bit numbers  $x, y$** 
  - Either signed or unsigned
- **But, exact results can be bigger than  $w$  bits**
  - Unsigned: up to  $2w$  bits
    - Result range:  $0 \leq x * y \leq (2^w - 1)^2 = 2^{2w} - 2^{w+1} + 1$
  - Two's complement min (negative): Up to  $2w-1$  bits
    - Result range:  $x * y \geq (-2^{w-1}) * (2^{w-1} - 1) = -2^{2w-2} + 2^{w-1}$
  - Two's complement max (positive): Up to  $2w$  bits, but only for  $(TMin_w)^2$ 
    - Result range:  $x * y \leq (-2^{w-1})^2 = 2^{2w-2}$
- **So, maintaining exact results...**
  - would need to keep expanding word size with each product computed
  - is done in software, if needed
    - e.g., by “arbitrary precision” arithmetic packages

# Unsigned Multiplication in C


Operands:  $w$  bits



True Product:  $2 \cdot w$  bits



Discard  $w$  bits:  $w$  bits



## ■ Standard Multiplication Function


- Ignores high order  $w$  bits

## ■ Implements Modular Arithmetic


$$\text{UMult}_w(u, v) = u \cdot v \bmod 2^w$$

# Signed Multiplication in C

Operands:  $w$  bits



True Product:  $2 \cdot w$  bits



Discard  $w$  bits:  $w$  bits



## ■ Standard Multiplication Function


- Ignores high order  $w$  bits
- Some of which are different for signed vs. unsigned multiplication
- Lower bits are the same

# Power-of-2 Multiply with Shift

## ■ Operation

- $u \ll k$  gives  $u * 2^k$
- Both signed and unsigned

Operands:  $w$  bits




## ■ Examples

- $u \ll 3 \quad \quad \quad == \quad u * 8$
- $(u \ll 5) - (u \ll 3) == \quad u * 24$
- Most machines shift and add faster than multiply
  - Compiler generates this code automatically

# Unsigned Power-of-2 Divide with Shift

## ■ Quotient of Unsigned by Power of 2

- $u \gg k$  gives  $\lfloor u / 2^k \rfloor$
- Uses logical shift



|                     | Division          | Computed     | Hex   | Binary            |
|---------------------|-------------------|--------------|-------|-------------------|
| <b>x</b>            | <b>15213</b>      | <b>15213</b> | 3B 6D | 00111011 01101101 |
| <b>x &gt;&gt; 1</b> | <b>7606.5</b>     | <b>7606</b>  | 1D B6 | 00011101 10110110 |
| <b>x &gt;&gt; 4</b> | <b>950.8125</b>   | <b>950</b>   | 03 B6 | 00000011 10110110 |
| <b>x &gt;&gt; 8</b> | <b>59.4257813</b> | <b>59</b>    | 00 3B | 00000000 00111011 |

# Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting
  - **Summary**
- Representations in memory, pointers, strings

# Arithmetic: Basic Rules

## ■ Addition:

- Unsigned/signed: Normal addition followed by truncate, same operation on bit level
- Unsigned: addition mod  $2^w$ 
  - Mathematical addition + possible subtraction of  $2^w$
- Signed: modified addition mod  $2^w$  (result in proper range)
  - Mathematical addition + possible addition or subtraction of  $2^w$

## ■ Multiplication:

- Unsigned/signed: Normal multiplication followed by truncate, same operation on bit level
- Unsigned: multiplication mod  $2^w$
- Signed: modified multiplication mod  $2^w$  (result in proper range)

# Why Should I Use Unsigned?

## ■ *Don't* use without understanding implications

- Easy to make mistakes

```
unsigned i;  
for (i = cnt-2; i >= 0; i--)  
    a[i] += a[i+1];
```

- Can be very subtle

```
#define DELTA sizeof(int)  
int i;  
for (i = CNT; i-DELTA >= 0; i-= DELTA)  
    . . .
```



# Counting Down with Unsigned

## ■ Proper way to use unsigned as loop index

```
unsigned i;  
for (i = cnt-2; i < cnt; i--)  
    a[i] += a[i+1];
```

## ■ See Robert Seacord, *Secure Coding in C and C++*

- C Standard guarantees that unsigned addition will behave like modular arithmetic
  - $0 - 1 \rightarrow UMax$

## ■ Even better

```
size_t i;  
for (i = cnt-2; i < cnt; i--)  
    a[i] += a[i+1];
```

- Data type `size_t` defined as unsigned value with length = word size
- Code will work even if `cnt = UMax`
- What if `cnt` is signed and `< 0`?


## Why Should I Use Unsigned? (cont.)

- ***Do Use When Performing Modular Arithmetic***
  - Multiprecision arithmetic
- ***Do Use When Using Bits to Represent Sets***
  - Logical right shift, no sign extension

# Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
  - Representation: unsigned and signed
  - Conversion, casting
  - Expanding, truncating
  - Addition, negation, multiplication, shifting
  - Summary
- **Representations in memory, pointers, strings**

# Byte-Oriented Memory Organization



- **Programs refer to data by address**
  - Conceptually, envision it as a very large array of bytes
    - In reality, it's not, but can think of it that way
  - An address is like an index into that array
    - and, a pointer variable stores an address
- **Note: system provides private address spaces to each “process”**
  - Think of a process as a program being executed
  - So, a program can clobber its own data, but not that of others


# Machine Words

- **Any given computer has a “Word Size”**
  - Nominal size of integer-valued data
    - and of addresses
  - Until recently, most machines used 32 bits (4 bytes) as word size
    - Limits addresses to 4GB ( $2^{32}$  bytes)
  - Increasingly, machines have 64-bit word size
    - Potentially, could have 18 EB (exabytes) of addressable memory
    - That's  $18.4 \times 10^{18}$
  - Machines still support multiple data formats
    - Fractions or multiples of word size
    - Always integral number of bytes

# Word-Oriented Memory Organization

## ■ Addresses Specify Byte Locations

- Address of first byte in word
- Addresses of successive words differ by 4 (32-bit) or 8 (64-bit)



# Example Data Representations

| C Data Type              | Typical 32-bit | Typical 64-bit | x86-64 |
|--------------------------|----------------|----------------|--------|
| <code>char</code>        | 1              | 1              | 1      |
| <code>short</code>       | 2              | 2              | 2      |
| <code>int</code>         | 4              | 4              | 4      |
| <code>long</code>        | 4              | 8              | 8      |
| <code>float</code>       | 4              | 4              | 4      |
| <code>double</code>      | 8              | 8              | 8      |
| <code>long double</code> | –              | –              | 10/16  |
| <code>pointer</code>     | 4              | 8              | 8      |

# Byte Ordering

- So, how are the bytes within a multi-byte word ordered in memory?
- Conventions
  - Big Endian: Sun, PPC Mac, Internet
    - Least significant byte has highest address
  - Little Endian: x86, ARM processors running Android, iOS, and Windows
    - Least significant byte has lowest address




# Byte Ordering Example


## ■ Example

- Variable x has 4-byte value of 0x01234567
- Address given by &x is 0x100

### Big Endian



### Little Endian




# Representing Integers

Decimal: 15213


Binary: 0011 1011 0110 1101

Hex: 3 B 6 D


`int A = 15213;`



`long int C = 15213;`



`int B = -15213;`



Two's complement representation

# Examining Data Representations

## ■ Code to Print Byte Representation of Data

- Casting pointer to unsigned char \* allows treatment as a byte array

```
typedef unsigned char *pointer;

void show_bytes(pointer start, size_t len) {
    size_t i;
    for (i = 0; i < len; i++)
        printf("%p\t0x%.2x\n", start+i, start[i]);
    printf("\n");
}
```

### Printf directives:

%p: Print pointer

%x: Print Hexadecimal

## show\_bytes Execution Example

```
int a = 15213;  
printf("int a = 15213;\n");  
show_bytes((pointer) &a, sizeof(int));
```

### Result (Linux x86-64):

```
int a = 15213;  
0x7fffb7f71dbc    6d  
0x7fffb7f71dbd    3b  
0x7fffb7f71dbe    00  
0x7fffb7f71dbf    00
```

# Representing Pointers

```
int B = -15213;  
int *P = &B;
```

**Sun**

|    |
|----|
| EF |
| FF |
| FB |
| 2C |

**IA32**

|    |
|----|
| AC |
| 28 |
| F5 |
| FF |

**x86-64**

|    |
|----|
| 3C |
| 1B |
| FE |
| 82 |
| FD |
| 7F |
| 00 |
| 00 |

**Different compilers & machines assign different locations to objects**

**Even get different results each time run program**

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

61

# Representing Strings


```
char S[6] = "18213";
```

## ■ Strings in C

- Represented by array of characters
- Each character encoded in ASCII format
  - Standard 7-bit encoding of character set
  - Character "0" has code 0x30
    - Digit  $i$  has code  $0x30+i$
- String should be null-terminated
  - Final character = 0

## ■ Compatibility

- Byte ordering not an issue



# Integer C Puzzles

## Initialization

```
int x = foo();
int y = bar();
unsigned ux = x;
unsigned uy = y;
```

- `x < 0`                      ☐ ☐ `((x*2) < 0)`
- `ux >= 0`
- `x & 7 == 7`                ☐ ☐ `(x<<30) < 0`
- `ux > -1`
- `x > y`                      ☐ ☐ `-x < -y`
- `x * x >= 0`
- `x > 0 && y > 0`        ☐ ☐ `x + y > 0`
- `x >= 0`                    ☐ ☐ `-x <= 0`
- `x <= 0`                    ☐ ☐ `-x >= 0`
- `(x|-x)>>31 == -1`
- `ux >> 3 == ux/8`
- `x >> 3 == x/8`
- `x & (x-1) != 0`


# Bonus extras



# Application of Boolean Algebra

## ■ Applied to Digital Systems by Claude Shannon

- 1937 MIT Master's Thesis
- Reason about networks of relay switches
  - Encode closed switch as 1, open switch as 0



Connection when

$$A \& \sim B \mid \sim A \& B$$

$$= A \wedge B$$

# Binary Number Property

## Claim

$$1 + 1 + 2 + 4 + 8 + \dots + 2^{w-1} = 2^w$$

$$1 + \sum_{i=0}^{w-1} 2^i = 2^w$$

- **w = 0:**

- $1 = 2^0$

- **Assume true for w-1:**

- $1 + 1 + 2 + 4 + 8 + \dots + 2^{w-1} + 2^w = 2^w + 2^w = 2^{w+1}$


$$= 2^w$$

# Code Security Example

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

- Similar to code found in FreeBSD's implementation of `getpeername`
- There are legions of smart people trying to find vulnerabilities in programs

# Typical Usage

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

```
#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, MSIZE);
    printf("%s\n", mybuf);
}
```

# Malicious Usage

```
/* Declaration of library function memcpy */
void *memcpy(void *dest, void *src, size_t n);
```

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

```
#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, -MSIZE);
    . . .
}
```

# Mathematical Properties

## ■ Modular Addition Forms an *Abelian Group*

- **Closed** under addition

$$0 \leq \text{UAdd}_w(u, v) \leq 2^w - 1$$

- **Commutative**

$$\text{UAdd}_w(u, v) = \text{UAdd}_w(v, u)$$

- **Associative**

$$\text{UAdd}_w(t, \text{UAdd}_w(u, v)) = \text{UAdd}_w(\text{UAdd}_w(t, u), v)$$

- **0** is additive identity

$$\text{UAdd}_w(u, 0) = u$$

- Every element has additive **inverse**

- Let  $\text{UComp}_w(u) = 2^w - u$   
 $\text{UAdd}_w(u, \text{UComp}_w(u)) = 0$

# Mathematical Properties of TAdd

## ■ Isomorphic Group to unsigneds with UAdd

- $TAdd_w(u, v) = U2T(UAdd_w(T2U(u), T2U(v)))$ 
  - Since both have identical bit patterns

## ■ Two's Complement Under TAdd Forms a Group

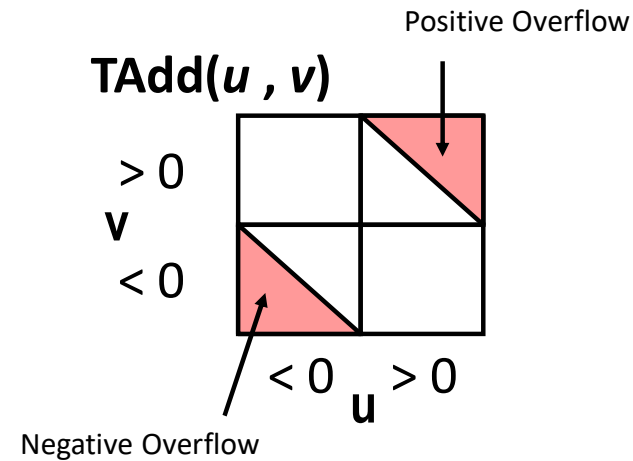
- Closed, Commutative, Associative, 0 is additive identity
- Every element has additive inverse

$$TComp_w(u) = \begin{cases} -u & u \neq TMin_w \\ TMin_w & u = TMin_w \end{cases}$$

# Characterizing TAdd

## ■ Functionality

- True sum requires  $w+1$  bits
- Drop off MSB
- Treat remaining bits as 2's comp. integer



$$TAdd_w(u, v) = \begin{cases} u + v + 2^w & u + v < TMin_w \text{ (NegOver)} \\ u + v & TMin_w \leq u + v \leq TMax_w \\ u + v - 2^w & TMax_w < u + v \text{ (PosOver)} \end{cases}$$



# Negation: Complement & Increment

## ■ Claim: Following Holds for 2's Complement

$$\sim x + 1 == -x$$

## ■ Complement

- Observation:  $\sim x + x == 1111\dots111 == -1$

$$\begin{array}{r}
 x \quad 10011101 \\
 + \quad \sim x \quad 01100010 \\
 \hline
 -1 \quad 11111111
 \end{array}$$

## ■ Complete Proof?

# Complement & Increment Examples

**x = 15213**

|             | Decimal       | Hex   | Binary            |
|-------------|---------------|-------|-------------------|
| <b>x</b>    | <b>15213</b>  | 3B 6D | 00111011 01101101 |
| <b>~x</b>   | <b>-15214</b> | C4 92 | 11000100 10010010 |
| <b>~x+1</b> | <b>-15213</b> | C4 93 | 11000100 10010011 |
| <b>y</b>    | <b>-15213</b> | C4 93 | 11000100 10010011 |

**x = 0**


|             | Decimal   | Hex   | Binary            |
|-------------|-----------|-------|-------------------|
| <b>0</b>    | <b>0</b>  | 00 00 | 00000000 00000000 |
| <b>~0</b>   | <b>-1</b> | FF FF | 11111111 11111111 |
| <b>~0+1</b> | <b>0</b>  | 00 00 | 00000000 00000000 |

## Code Security Example #2

### ■ SUN XDR library

- Widely used library for transferring data between machines

```
void* copy_elements(void *ele_src[], int ele_cnt, size_t ele_size);
```



## XDR Code

```
void* copy_elements(void *ele_src[], int ele_cnt, size_t ele_size) {
    /*
     * Allocate buffer for ele_cnt objects, each of ele_size bytes
     * and copy from locations designated by ele_src
     */
    void *result = malloc(ele_cnt * ele_size);
    if (result == NULL)
        /* malloc failed */
        return NULL;
    void *next = result;
    int i;
    for (i = 0; i < ele_cnt; i++) {
        /* Copy object i to destination */
        memcpy(next, ele_src[i], ele_size);
        /* Move pointer to next memory region */
        next += ele_size;
    }
    return result;
}
```

# XDR Vulnerability

`malloc(ele_cnt * ele_size)`

- What if:

- `ele_cnt` =  $2^{20} + 1$
- `ele_size` = 4096 =  $2^{12}$
- Allocation = ??

- How can I make this function secure?

# Compiled Multiplication Code

## C Function

```
long mul12(long x)
{
    return x*12;
}
```

## Compiled Arithmetic Operations

```
leaq (%rax,%rax,2), %rax
salq $2, %rax
```

## Explanation

```
t <- x+x*2
return t << 2;
```

- C compiler automatically generates shift/add code when multiplying by constant

# Compiled Unsigned Division Code

## C Function

```
unsigned long udiv8
(unsigned long x)
{
    return x/8;
}
```

## Compiled Arithmetic Operations

```
shrq $3, %rax
```

## Explanation


```
# Logical shift
return x >> 3;
```

- Uses logical shift for unsigned
- For Java Users
  - Logical shift written as >>>

# Signed Power-of-2 Divide with Shift

## ■ Quotient of Signed by Power of 2

- $x \gg k$  gives  $\lfloor x / 2^k \rfloor$
- Uses arithmetic shift
- Rounds wrong direction when  $u < 0$



|           | Division    | Computed | Hex   | Binary            |
|-----------|-------------|----------|-------|-------------------|
| $y$       | -15213      | -15213   | C4 93 | 11000100 10010011 |
| $y \gg 1$ | -7606.5     | -7607    | E2 49 | 11100010 01001001 |
| $y \gg 4$ | -950.8125   | -951     | FC 49 | 11111100 01001001 |
| $y \gg 8$ | -59.4257813 | -60      | FF C4 | 11111111 11000100 |




# Correct Power-of-2 Divide

## ■ Quotient of Negative Number by Power of 2

- Want  $\lceil x / 2^k \rceil$  (Round Toward 0)
- Compute as  $\lfloor (x + 2^k - 1) / 2^k \rfloor$ 
  - In C:  $(x + (1 \ll k) - 1) \gg k$
  - Biases dividend toward 0

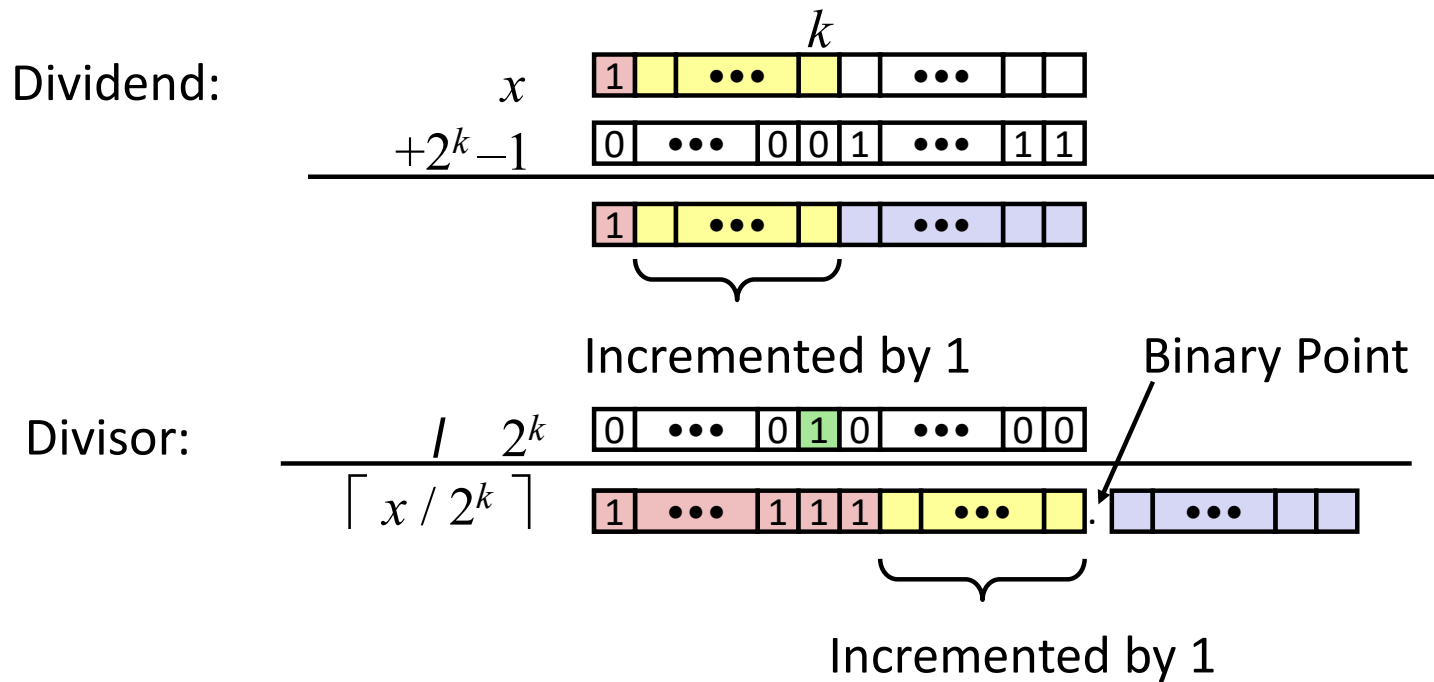
### Case 1: No rounding



***Biasing has no effect***

# Correct Power-of-2 Divide (Cont.)

## Case 2: Rounding



***Biasing adds 1 to final result***

# Compiled Signed Division Code

## C Function

```
long idiv8(long x)
{
    return x/8;
}
```

## Compiled Arithmetic Operations

```
    testq %rax, %rax
    js    L4
L3:
    sarq $3, %rax
    ret
L4:
    addq $7, %rax
    jmp  L3
```

## Explanation

```
if x < 0
    x += 7;
# Arithmetic shift
return x >> 3;
```

- Uses arithmetic shift for int
- For Java Users
  - Arith. shift written as >>

# Arithmetic: Basic Rules

- **Unsigned ints, 2's complement ints are isomorphic rings:  
isomorphism = casting**
  
- **Left shift**
  - Unsigned/signed: multiplication by  $2^k$
  - Always logical shift
  
- **Right shift**
  - Unsigned: logical shift, div (division + round to zero) by  $2^k$
  - Signed: arithmetic shift
    - Positive numbers: div (division + round to zero) by  $2^k$
    - Negative numbers: div (division + round away from zero) by  $2^k$   
Use biasing to fix

# Properties of Unsigned Arithmetic

## ■ Unsigned Multiplication with Addition Forms Commutative Ring

- Addition is commutative group
- Closed under multiplication
$$0 \leq \text{UMult}_w(u, v) \leq 2^w - 1$$
- Multiplication Commutative
$$\text{UMult}_w(u, v) = \text{UMult}_w(v, u)$$
- Multiplication is Associative
$$\text{UMult}_w(t, \text{UMult}_w(u, v)) = \text{UMult}_w(\text{UMult}_w(t, u), v)$$
- 1 is multiplicative identity
$$\text{UMult}_w(u, 1) = u$$
- Multiplication distributes over addition
$$\text{UMult}_w(t, \text{UAdd}_w(u, v)) = \text{UAdd}_w(\text{UMult}_w(t, u), \text{UMult}_w(t, v))$$

# Properties of Two's Comp. Arithmetic

## ■ Isomorphic Algebras

- Unsigned multiplication and addition
  - Truncating to  $w$  bits
- Two's complement multiplication and addition
  - Truncating to  $w$  bits

## ■ Both Form Rings

- Isomorphic to ring of integers mod  $2^w$

## ■ Comparison to (Mathematical) Integer Arithmetic

- Both are rings
- Integers obey ordering properties, e.g.,

$$u > 0 \quad \Rightarrow \quad u + v > v$$

$$u > 0, v > 0 \quad \Rightarrow \quad u \cdot v > 0$$

- These properties are not obeyed by two's comp. arithmetic

$$TMax + 1 == TMin$$

$$15213 * 30426 == -10030 \quad (16\text{-bit words})$$

# Reading Byte-Reversed Listings

## ■ Disassembly

- Text representation of binary machine code
- Generated by program that reads the machine code

## ■ Example Fragment

| Address  | Instruction Code            | Assembly Rendition    |
|----------|-----------------------------|-----------------------|
| 8048365: | 5b                          | pop %ebx              |
| 8048366: | 81 c3 <u>ab 12 00 00</u>    | add \$0x12ab,%ebx     |
| 804836c: | 83 bb 28 00 <u>00 00 00</u> | cmpl \$0x0,0x28(%ebx) |

## ■ Deciphering Numbers

- Value: 0x12ab
- Pad to 32 bits: 0x000012ab
- Split into bytes: 00 00 12 ab
- Reverse: ab 12 00 00